

Solving the TTC Families to Persons Case with FunnyQT

Tassilo Horn
tsdh@gnu.org

The GNU Project

Abstract

This paper describes the FunnyQT solution to the bidirectional TTC 2017 Families to Persons transformation case. The solution is simple and concise and passes all batch transformation and some incremental tests.

1 Introduction

This paper describes the FunnyQT¹ [Hor16, Hor15] solution of the TTC 2017 Families to Persons case [ABW17]. With only 52 lines of declarative code, the solution is able to pass all batch transformation tests and some incremental tests. The solution project is available on Github² and it is integrated into the *benchmarkx*³ framework.

FunnyQT is a model querying and transformation library for the functional Lisp dialect Clojure⁴. Queries and transformations are Clojure functions using the features provided by the FunnyQT API.

Clojure provides strong metaprogramming capabilities that are used by FunnyQT in order to define several *embedded domain-specific languages* (embedded DSL[Fow10]) for different querying and transformation tasks.

FunnyQT is designed to be extensible. By default, it supports EMF [SBPM08] and JGraLab⁵ TGraph models. Support for other modeling frameworks can be added without having to touch FunnyQT's internals.

The FunnyQT API is structured into several namespaces, each namespace providing constructs supporting concrete querying and transformation use-cases, e.g., model management, functional querying, polymorphic functions, relational querying, pattern matching, in-place transformations, out-place transformations, bidirectional transformations, and some more. For solving the families to persons case, its bidirectional transformation and relational model querying DSLs have been used.

2 Solution Description

This section explains the FunnyQT solution. First, section 2.1 explains the actual transformation solving the case. Then, section 2.2 explains how the solution is integrated into the *benchmarkx* framework.

2.1 The Families to Persons Transformation

As a first step, relational querying APIs for the two metamodels are generated.

```
1 (rel/generate-metamodel-relations "metamodels/Families.ecore" f)
2 (rel/generate-metamodel-relations "metamodels/Persons.ecore" p)
```

This makes the relations for the families metamodel available with the namespace alias `f` and those of the persons metamodel with alias `p`. Relations are generated for every metamodel element. For example, there's a relation `(f/Family f family)` which succeeds for every `family` in the given family model `f`. There's a relation

Copyright © by the paper's authors. Copying permitted for private and academic purposes.

In: A. Editor, B. Coeditor (eds.): Proceedings of the XYZ Workshop, Location, Country, DD-MMM-YYYY, published at <http://ceur-ws.org>

¹<http://funnyqt.org> ²<https://github.com/tsdh/ttc17-families2persons-bx> ³<https://github.com/eMoflon/benchmarkx/>
⁴<http://clojure.org> ⁵<https://github.com/jgralab/jgralab>

(`f/name f el n`) which succeeds if `el` is a member or family of the family model `f` whose name is `n`. And there're reference relations like (`f/->sons f parent child`) which succeeds if `child` is a son of `parent`.

Before the transformation itself, we define a helper relation which defines the possible kinds of relationships between a `family` and a `family member` depending on if we prefer to create parents over creating children (parameter `pref-parent`). This is a higher-order relation in that the two remaining parameters are a parent relation `prel` (either `f/->father` or `f/->mother`) and a child relation (either `f/->daughters` or `f/->sons`).

```

3 (defn relationshipo [pref-parent f family member prel crel]
4   (ccl/conda
5     [(bx/target-directiono :right)      ;; (1)
6       (ccl/conde
7         [(prel f family member)]
8         [(crel f family member)])])
9     [(bx/existing-elemento? member)]    ;; (2)
10    [(ccl/= pref-parent false)           ;; (3)
11      (crel f family member)]
12    [(bx/unseto? f family prel member) ;; (4)
13      (prel f family member)]
14    [(crel f family member)])           ;; (5)

```

`ccl/conda` is like a short-cutting logical disjunction. The n -th clause is only tried if all preceding clauses fail. The first clause succeeds when we are transforming into the direction of the right model, i.e., the person register. In this case, `member` may be in a parental role of family (`prel`), or it might be in a child role (`crel`). We don't really care but want to ensure that all members of the given `family` are reachable, thus we use a non-short-cutting `ccl/conde` which gives both clauses a chance to succeed.

All other clauses deal with transforming in the direction of the family model, i.e., the backward transformation. The second clause deals with the case where `member` is an already existing element, i.e., not coming into being by the current execution of the transformation. Here, we assume that this member is already properly assigned to a family, so we simply succeed without doing anything.

In clause three, if we do not prefer assigning to parental roles, then the child relation `crel` must succeed between the `family` and the `member`. The child relations can always succeed because a family can have an arbitrary number of daughters and sons. Thus, in the remaining clauses we have to deal only with the case that we do prefer assigning to parental roles.

In the fourth clause, if the `family`'s parental role is still unset or already assigned to `member`, then the parental relation must succeed between the `family` and the `member`.

Lastly, if no clause has succeeded until now, then the child relation has to succeed. As said, this goal⁶ can always succeed.

In the following, the actual transformation definition is explained. It is defined using the `deftransformation` macro provided by FunnyQT.

```

15 (bx/deftransformation families2persons [f p prefer-parent prefer-ex-family]
16   :delete-unmatched-target-elements true
17   :id-init-fn bx/number-all-source-model-elements

```

The transformation's name is `families2persons` and it declares four parameters. The parameter `f` is the family model (the left model), `p` is the persons model (the right model), `prefer-parent` is a boolean flag determining if we prefer creating parents to creating children, and `prefer-ex-family` is a boolean flag, too, determining if we prefer re-using existing families over creating new families for new family members.

By default, bidirectional FunnyQT transformations will never delete elements from the current target model, i.e., the model in whose direction the synchronization is performed. The reason for that behavior is that it allows to run the transformation in one direction first and then in the other direction in order to perform a full synchronization where missing elements are created in each of the two models. Thus, after running a transformation, e.g., in the direction of the right model, it is only ensured that for each element (considered by the transformation's rules) in the left model, there is a corresponding counterpart in the right model. However, the right model might still contain elements which have no counterpart in the left model. With option `:delete-unmatched-target-elements` set to `true`, this behavior is changed. Elements in the current target model which are not required by the current source model and the transformation relations are deleted.

The next option, `:id-init-fn`, has the following purpose. In this transformation case, family members and persons don't have some kind of unique identity. For example, it is allowed to have two members named Jim with

⁶The application of a relation is called a *goal*.

the same name in the very same family Smith. However, FunnyQT BX transformations have forall-there-exists semantics where it would suffice to create just one person in the right model with the name set to “Smith, Jim”. But the case description mandates that we create one person for every member and vice versa, no matter if they can be distinguished based on their attribute values. For such scenarios, FunnyQT’s bidirectional transformation DSL provides a concept of synthetic ID attributes. The value of `:id-init-fn` has to be a function which returns a map from elements to their synthetic IDs. The built-in function `bx/number-all-source-model-elements` returns a map where every element in the source model gets assigned a unique integer number. These synthetic IDs are then used in a transformation relation which is discussed further below.

The first transformation relation, `family-register2person-register`, transforms between family and person registers.

```

18  (~:top family-register2person-register
19  :left  [(f/FamilyRegister f ?family-register)]
20  :right [(p/PersonRegister p ?person-register)]
21  :where [(member2female :?family-register ?family-register :?person-register ?person-register)
22           (member2male :?family-register ?family-register :?person-register ?person-register)]

```

It is defined as a top-level rule meaning that it’ll be executed as the transformation’s entry point. It’s `:left` and `:right` clauses describe that for every `?family-register` there has to be a `?person-register` and vice versa. We assume that there’s always just one register in each model.

The `:where` clause defines that after this relation has been enforced (or checked in checkonly mode), then the two transformation relations `member2female` and `member2male` have to be enforced (or tested) between the current `?family-register` and `?person-register`⁷.

The next transformation relation, `member2person`, describes how family members of a family contained in a family register in the left model correspond to persons contained in a person register in the right model. As can be seen, there’s no goal describing how the `?family` and the `?member` are connected in the `:left` clause and in the `:right` clause we’re dealing with just a `?person` of type `Person` which is an abstract class. As such, this relation is not sufficient for the complete synchronization between members in the different roles of a family to females and males. Instead, it only captures the aspects that are common in the cases where mothers and daughters are synchronized with females and fathers and sons are synchronized with males. Therefore, this relation is declared abstract.

```

23  (~:abstract member2person
24  :left  [(f/->families f ?family-register ?family)
25          (f/Family f ?family)
26          (f/name f ?family ?last-name)
27          (f/FamilyMember f ?member)
28          (f/name f ?member ?first-name)
29          (id ?member ?id)
30          (ccl/conda
31            [(ccl== prefer-ex-family true)]
32            [(bx/existing-elemento? ?member)
33              (id ?family ?last-name)]
34            [(id ?family ?id)]]
35  :right [(p/->persons p ?person-register ?person)
36          (p/Person p ?person)
37          (p/name p ?person ?full-name)
38          (id ?person ?id)]
39  :when  [(rel/stro ?last-name ", " ?first-name ?full-name))]

```

So what are these common aspects? Well, a `?member` of a `?family` (where we haven’t determined the role, yet) contained in the `?family-register` passed in as parameter from `family-register2person-register` corresponds to a `?person` (where we haven’t determined the gender, yet) contained in the `?person-register` passed in as the other parameter from `family-register2person-register`. The `:when` clause defines that the concatenation of the `?family`’s `?last-name`, the string `", "` and the `?member`’s `?first-name` gives the `?full-name` of the `?person`.

What hasn’t been described so far are the `id` goals in lines 29, and 34 and the `ccl/conda` goal starting in line 30. The first two define that the `?member` and the corresponding `?person` must have the same synthetic ID. Remember the `:id-init-fn` in line 17 which assigned a unique number to every element in the respective source model of the transformation. With these synthetic IDs, the transformation is able to create one person for every member and vice versa even in the case where two elements are equal based on attribute values.

⁷Transformation relations are called with keyword parameters. The two calls in the `:where` clause say that the current `?family-register` will be bound to the logic variable with the same name in the called relation and the same is true for the `?person-register`.

Lastly, the `ccl/conda` goal starting in line 30 of the `:left` clause handles the preference of re-using existing families, i.e., assigning new members to existing families, over creating new families for new members. By default, FunnyQT would always try to re-use an existing family. Thus, if the `prefer-ex-family` parameter is `true`, nothing needs to be done. Likewise, if `?member` is an existing element for which we assume she's already assigned to some family, we can also just stick to the default behavior but define the `?family`'s ID to be its name (although it's probably not unique). If the first two `ccl/conda` clauses fail, i.e., `prefer-ex-family` is `false` and `?member` is a new member which is just going to be created by the enforcement of this relation, then we define that the `?family`'s ID must equal the IDs of the `?member` and `?person`. Thus, in this case and only in this case, new members force the creation of a new family even when there already is a family with the right name.

The last two transformation relations extend the `member2person` relation for synchronizing between members in the role of a family mother or daughter and female persons and between family fathers or sons and male persons.

```

40 (member2female
41  :extends [(member2person)]
42  :left [(relationshipo prefer-parent f ?family ?member f/->mother f/->daughters)]
43  :right [(p/Female p ?person)])
44 (member2male
45  :extends [(member2person)]
46  :left [(relationshipo prefer-parent f ?family ?member f/->father f/->sons)]
47  :right [(p/Male p ?person)])

```

In the `:left` clauses we use the `relationshipo` helper relation described in the beginning of this section which chooses the right female or male role based on the preference parameter `prefer-parent` and the current state of the family, i.e., by checking if the respective parental role is still unset. In the two `:right` clauses, we only need to specify that the `Person ?person` is actually a `Female` or `Male`.

These 33 lines of transformation specification plus the 12 lines for the `relationshipo` helper, and two lines for the generation of the metamodel-specific relational querying APIs form the complete functional parts of the solution. The only thing omitted from the paper are the namespace declaration⁸ consisting of 5 lines of code.

2.2 Gluing the Solution with the Framework

Typically, open-source Clojure libraries and programs are distributed as JAR files that contain the source files rather than byte-compiled class files. This solution does almost the same except that the JAR contains the solution source code, FunnyQT itself (also as sources) and every dependency of FunnyQT (like Clojure) except for EMF which the *benchmark* project already provides.

Calling Clojure functions from Java is really easy and FunnyQT transformations are no exception because they are plain Clojure functions, too. The FunnyQT solution's `BXTool` implementation `FunnyQTFamiliesToPerson` extends the `BXToolForEMF` class. Essentially, it just has a static member `T` which is set to the transformation.

```

public class FunnyQTFamiliesToPerson extends BXToolForEMF<FamilyRegister, PersonRegister, Decisions> {
    private final static Keyword LEFT = (Keyword) Clojure.read(":left");
    private final static Keyword RIGHT = (Keyword) Clojure.read(":right");
    private final static IFn T; // <-- The FunnyQT Transformation

    static {
        final String transformationNamespace = "ttc17-families2persons-bx.core";
        // Clojure's require is similar to Java's import. However, it also loads the required
        // namespace from a source code file and immediately compiles it.
        final IFn require = Clojure.var("clojure.core", "require");
        require.invoke(Clojure.read(transformationNamespace));
        T = Clojure.var(transformationNamespace, "families2persons");
    }
}

```

All Clojure functions implement the `IFn` interface and can be called using `invoke()`. And exactly this is done to call the transformation.

```

private void transform(Keyword direction) {
    T.invoke(srcModel, trgModel, direction,
            configurator.decide(Decisions.PREFER_CREATING_PARENT_TO_CHILD),
            configurator.decide(Decisions.PREFER_EXISTING_FAMILY_TO_NEW));
}

```

This corresponds to a direct Clojure call (`families2persons src trg dir prefer-parent prefer-ex-family`).

⁸The Clojure equivalent of Java's package statement and imports.

3 Evaluation & Conclusion

In this section, the FunnyQT solution’s test results are presented and classified as requested by the case description. Since FunnyQT’s bidirectional transformation DSL is state-based and not incremental at all (and by design), many of the incremental tests are mostly out of scope. However, in its core use case, non-incremental bidirectional transformations, the solution passes all tests.

BatchForward.* All tests result in an *expected pass*.

BatchBwdEandP.* All tests result in an *expected pass*.

BatchBwdEnotP.* All tests result in an *expected pass*.

BatchBwdNotEandP.* All tests result in an *expected pass*.

BatchBwdNotEnotP.* All tests result in an *expected pass*.

IncrementalForward.* The solution *expectedly passes* the tests `testStability`, `testHippocraticness`, and `testIncrementalInserts`. All remaining tests *fail expectedly* because those tests require incremental abilities. For example, in some tests the birthdates are set manually in the persons model. The re-execution of the transformation causes deletion and re-creation of those persons, however then they get assigned default birthdates and not the manually edited ones as the transformation doesn’t consider this attribute at all.

IncrementalBackward.* The solution *expectedly passes* the tests `testStability`, `testIncrementalInsertsFixedConfig`, `testIncrementalOperational`, and `testHippocraticness`. The test `testIncrementalInsertsDynamicConfig` *fails unexpectedly*. When we neither prefer existing families nor assigning to parental roles, the transformation still adds the new Seymore to an existing family in a parental role. All other tests are *expected fails* due to the fact that they require incremental capabilities.

In summary, the correctness is satisfying. There is only one test which fails unexpectedly. All other fails are expected and can hardly be solved by a non-incremental approach.

A very weak point of the solution is its performance. It is at least an order of magnitude slower than the other solutions already integrated in the benchmarx project (BiGUL, eMoflon, BXTend, MediniQVT). Where their runtimes are in the tenth of seconds or below, the FunnyQT solution takes seconds. With models in the size of thousands of elements, you might have to wait a bit for the transformation to finish. Only reason is that the FunnyQT BX implementation is state-based and thus needs to recompute the correspondences between all elements in the source and target models on every invocation instead of just propagating small deltas like incremental approaches do. Furthermore, FunnyQT’s bidirectional transformation DSL is built upon the relational programming library *core.logic*⁹ which is not tuned for performance but for simplicity and extensibility of its implementation and probably FunnyQT doesn’t use it in the best possible way.

Other good points of the solution are its conciseness and simplicity. With only 52 lines of code, it is by far the most concise solution currently integrated in the benchmarx project. And it is quite simple to understand. The only complexities it has arise from the different alternatives depending on external parameters.

References

- [ABW17] Anthony Anjorin, Thomas Buchmann, and Bernhard Westfechtel. The family to persons case. In *Transformation Tool Contest*, 2017.
- [Fow10] Martin Fowler. *Domain-Specific Languages*. Addison-Wesley Professional, 2010.
- [Hor15] Tassilo Horn. Graph pattern matching as an embedded clojure dsl. In *International Conference on Graph Transformation - 8th International Conference, ICGT 2015, L’Aquila, Italy, July 2015*, 2015.
- [Hor16] Tassilo Horn. *A Functional, Comprehensive and Extensible Multi-Platform Querying and Transformation Approach*. PhD thesis, University Koblenz-Landau, Germany, 2016.
- [SBPM08] Dave Steinberg, Frank Budinsky, Marcelo Paternostro, and Ed Merks. *EMF: Eclipse Modeling Framework*. Addison-Wesley Professional, 2 edition, 2008.

⁹<https://github.com/clojure/core.logic>